

AI

Cloud Network Security Advisor

AI-powered security posture analysis · gap detection
hardening recommendations · 90-day action plan

Demo Mode

Claude API

OpenAI API

danvzla.github.io/ai-security-advisor

Daniel Mazzini · Principal Architect & Senior TPM



THE PROBLEM

Security posture analysis is slow, manual, and starts from scratch every time



Days of manual work

Full gap analysis and hardening recommendations take 2–3 days of senior architect time. Every engagement starts from scratch.



No structured output

Analysis lives in slide decks and email threads. No standardized risk scores, no sortable gap register, no client-ready action plan.



Tool expertise required

VMware NSX, SASE, Zero Trust, and compliance framework recommendations require deep specialization few architects have consistently.

THREE MODES

Zero setup to live AI analysis — one click apart



No API key required

Demo Mode

Click Demo → select a scenario → click Analyze. Full pre-built analysis loads in under 1 second. All 4 tabs populate. Works offline, works anywhere.

Best for: Always ready — no setup, no key, no network dependency



claude-haiku-4-5

Claude API

Switch to Claude → enter sk-ant-... key → Connect → Analyze. Real Anthropic API call.

Recommendations grounded in actual LLM reasoning for your inputs.



gpt-4o-mini

OpenAI API

Switch to OpenAI → enter sk-... key → Connect → Analyze. Identical prompt, identical output structure, identical 4 tabs. One click to switch.

Everything a security engagement produces — in one analysis

01

Overview

Risk score (0–100), overall risk level, executive summary, and 4 key metrics. The CFO view of the security posture.

- ◆ Risk score with color-coded severity ring
- ◆ Critical / High / Medium gap counts
- ◆ Executive summary paragraph
- ◆ Quick-reference metric cards

02

Security Gaps

Every gap identified — sorted by severity. Critical first. Each card shows the problem, why it matters, and the specific fix.

- ◆ Sorted Critical → High → Medium → Low
- ◆ Color-coded severity pill per gap
- ◆ Root cause description
- ◆ Specific actionable recommendation

03

Hardening Recs

Prioritized hardening actions with tool alignment. Immediate actions clearly separated from short-term and medium-term.

- ◆ Immediate / Short-term / Medium-term priority
- ◆ What to do and exactly how
- ◆ Tool alignment (NSX, SASE, Vault, WAF, etc.)
- ◆ Color-coded border by urgency

04

90-Day Plan

Three structured delivery phases with timelines — ready to share with the client as a formal remediation roadmap.

- ◆ Phase 1: Critical remediation — Week 1-2
- ◆ Phase 2: Architecture hardening — Week 3-6
- ◆ Phase 3: Zero Trust maturity — Week 7-12
- ◆ Teal phase number indicators

HOW IT WORKS

Four steps · from environment description to full analysis

1

0 seconds

Select mode

Choose Demo (no key), Claude, or OpenAI. In Demo Mode — click Analyze immediately. No setup, no dependencies.

2

30 seconds

Describe environment

Enter network topology, select security tools from 12 options, compliance frameworks from 8 options, environment type and maturity level.

3

1–15 seconds

Run analysis

Click Analyze Security Posture. Demo: pre-built output in 1 second. API Mode: real LLM call tailored to your specific inputs.

4

Immediately

4 tabs of output

Overview · Security Gaps · Hardening Recommendations · 90-Day Action Plan. Export to Excel for client delivery.

Three layers · browser-native · no server required

PRESENTATION

- ◆ `index.html` — nav, mode bar, form, results container
- ◆ `styles.css` — dark slate design system (#0F172A)
- ◆ Mode selector: Demo / Claude / OpenAI pills
- ◆ Form: topology textarea, chip selectors, selects
- ◆ Results: 4-tab panel, `switchTab()` renderer
- ◆ Sticky header · standalone file · zero dependencies

GitHub Pages ready
Zero infrastructure cost

ORCHESTRATION

- ◆ `setMode()` — switches provider, updates UI status
- ◆ `runAnalysis()` — routes to demo or API path
- ◆ `buildPrompt()` — assembles structured LLM prompt
- ◆ `switchTab()` — tab panel controller
- ◆ `renderResults()` — populates all 4 output tabs
- ◆ `exportToExcel()` — multi-sheet xlsx download

Pure JavaScript
No build step required


AI + DATA

- ◆ `DEMO_ANALYSIS` — pre-built Telco/Hybrid output
- ◆ `callClaude()` — Anthropic API · haiku-4-5
- ◆ `callOpenAI()` — OpenAI API · gpt-4o-mini
- ◆ `parseJSON()` — 4-attempt JSON repair chain
- ◆ `SAMPLE` — pre-filled scenario topology
- ◆ `COMPLIANCE_VALS` — framework validation list

Claude or OpenAI
1-click provider swap

Data flow · JSON repair · provider abstraction

INPUT → PROCESSING → OUTPUT



JSON Repair Chain

LLMs sometimes return malformed JSON — especially near token limits. The 4-attempt repair chain ensures resilience:

1 **Direct JSON.parse()**

2 **Strip control chars**

3 **Collapse newlines**

4 **Close open brackets**

Attempt 4 handles token-limit truncation by closing all unclosed { and [brackets

SAMPLE OUTPUT

Hybrid cloud · NIST CSF + PCI-DSS · Defined maturity

71

RISK

3

Critical

5

High

4

Medium

9

Recs

" The hybrid multi-cloud environment has critical gaps in east-west segmentation, MFA enforcement, and secret management. Three critical findings require immediate remediation. NIST CSF and PCI-DSS compliance is at risk without structured action in the next 30 days. "

Critical

No micro-segmentation on east-west traffic

→ Deploy NSX distributed firewall with deny-all default. Map flows with NSX Intelligence.

Critical

MFA not enforced on privileged console access

→ Enforce FIDO2/hardware-key MFA on all privileged accounts within 72 hours.

High

No WAF on public-facing API endpoints

→ Deploy AWS WAF or Azure Front Door with managed OWASP rule sets. Enable rate limiting.

High

Secrets hardcoded in CI/CD pipeline scripts

→ Migrate all secrets to HashiCorp Vault or AWS Secrets Manager. Rotate all exposed keys.

KEY BENEFITS

What this demonstrates — and what it replaces

< 1s

Demo Mode

Full analysis, no API key

4

Output tabs

Overview · Gaps · Hardening · Plan

3

Providers

Demo · Claude API · OpenAI API

12+

Gaps detected

Sorted by severity, with fixes

9

Hardening recs

Immediate · Short-term · Medium-term

xlsx

Excel export

Multi-sheet client deliverable

SAMPLE SCENARIOS

Three pre-built scenarios — ready in Demo Mode

01

VMware VCF Private Cloud

Tools: VMware NSX · vDefend ATP · SD-WAN · SASE

Compliance: NIST CSF · SOC 2 · Maturity: Defined

Best for

VMware / Broadcom · Cisco · adjacent SIs

02

Hybrid Cloud (AWS + Azure)

Tools: Traditional Firewall · SIEM · EDR/XDR · IDS/IPS

Compliance: PCI-DSS · ISO 27001 · Maturity: Developing

Best for

Cognizant · Accenture · cloud SI accounts

03

Telco / CSP Network

Tools: SD-WAN · SASE · Zero Trust · PAM

Compliance: HIPAA · FedRAMP · GDPR · Maturity: Managed

Best for

Nokia · Ericsson · AT&T · T-Mobile · Verizon

◆ All three scenarios run in Demo Mode — no API key, no network dependency, always ready

LIVE DEMO · NO API KEY REQUIRED

Try it right now

[Live demo → danvzla.github.io/ai-security-advisor](https://danvzla.github.io/ai-security-advisor)

[GitHub → github.com/danvzla/ai-security-advisor](https://github.com/danvzla/ai-security-advisor)

- ◆ Open URL · click Demo Mode · click Analyze · full analysis in under 1 second, no key
- ◆ 4 output tabs: Overview · Security Gaps · Hardening Recommendations · 90-Day Plan
- ◆ 3 pre-built scenarios: VCF Private Cloud · Hybrid Cloud · Telco/CSP
- ◆ Switch to Claude or OpenAI mode for live AI generation with your own key

Daniel Mazzini

Principal Architect & Senior TPM · Cloud & Infrastructure · Telco · Applied AI

AI · Security Advisor